

# *W4111 – Introduction to Databases*

## *Section 002, Spring 2026*

### *Lecture 2: Foundation – ER, Relational, SQL (I)*



# Contents

- Show HW tips and guidelines on Ed.

# Contents

---

- ER Model/Modeling:
  - More details
  - An example.
- A deeper dive on *keys*.
  - Formal model for keys.
  - Some practical concepts for keys.
- Relational Algebra:
  - $\pi$ ,  $\sigma$  examples
  - $\times$  (Cartesian Product)
  - $\bowtie$  (JOIN) operator
- Continuing SQL
  - More details on SELECT
  - Cartesian Product
  - Join

# Entity Relationship Entity Relationship Models


# Modeling (Reminder/Review)

# Modeling According to Wikipedia

- Model-driven engineering (MDE) is a software development methodology that focuses on creating and exploiting domain models, which are conceptual models of all the topics related to a specific problem. Hence, it highlights and aims at abstract representations of the knowledge and activities that govern a particular application domain, rather than the computing (i.e. algorithmic) concepts.
- In software engineering, a domain model is a conceptual model of the domain that incorporates both behavior and data.
- A modeling language is any artificial language that can be used express data, information or knowledge or systems in a structure that is defined by a consistent set of rules. A modeling language can be graphical or textual.
  - Graphical modeling languages use a diagram technique with named symbols that represent concepts and lines that connect the symbols and represent relationships and various other graphical notation to represent constraints.
  - Textual modeling languages may use standardized keywords accompanied by parameters or natural language terms and phrases to make computer-interpretable expressions.

# ER Model and ER Modeling

- ER Model: Agility, Separation of Concerns
  - ER model is a generalization that most DB models implement in some form.
  - Using the ER model enables:
    - Thinking about and collaborating on design with getting bogged down in details.
    - Enable flexible choices about how to realize/Implement data.
- ER Diagrams: Communication, Quality, Precision
  - With a little experience, everyone can understand and ER diagram.
  - Easier to discuss and collaborate on application's data than showing SQL table definitions, JSON, ... ....
  - People think visually. That is why we have whiteboards. ER diagrams are precise and unambiguous.
  - Guides you to think about relationships, keys, ... And prevents “re-dos” later in the process. It is easier to fix a diagram than a database schema.



# ER Modeling – Reasonably Good Summary

## Advantages of ER Model

**Conceptually it is very simple:** ER model is very simple because if we know relationship between entities and attributes, then we can easily draw an ER diagram.

**Better visual representation:** ER model is a diagrammatic representation of any logical structure of database. By seeing ER diagram, we can easily understand relationship among entities and relationship.

**Effective communication tool:** It is an effective communication tool for database designer.

**Highly integrated with relational model:** ER model can be easily converted into relational model by simply converting ER model into tables.

**Easy conversion to any data model:** ER model can be easily converted into another data model like hierarchical data model, network data model and so on.

## Disadvantages of ER Model

**Limited constraints and specification**

**Loss of information content:** Some information be lost or hidden in ER model

**Limited relationship representation:** ER model represents limited relationship as compared to another data models like relational model etc.

**No representation of data manipulation:** It is difficult to show data manipulation in ER model.

**Popular for high level design:** ER model is very popular for designing high level design

**No industry standard for notation**

<https://pctechnicalpro.blogspot.com/2017/04/advantages-disadvantages-er-model-dbms.html>


### Note:

- If you get to use Google to help with take home exams, HW, etc.
- I get to use Google to help with slides.

# A Common and my Approach: Conceptual → Logical → Physical

<https://ehkioya.com/conceptual-logical-physical-database-modeling/>

## Levels of Data Modeling



- It is easy to get carried away with modeling. You can spend all your time modeling and not actually build the schema.
- We will use the approaches in class.
- Mostly to understand concepts and patterns.


<https://www.1keydata.com/datawarehousing/data-modeling-levels.html>

Feature	Conceptual	Logical	Physical
Entity Names	✓	✓	
Entity Relationships	✓	✓	
Attributes		✓	
Primary Keys		✓	✓
Foreign Keys		✓	✓
Table Names			✓
Column Names			✓
Column Data Types			✓

Conceptual Model Design


Logical Model Design

Physical Model Design



<https://www.1keydata.com/datawarehousing/data-modeling-levels.html>

# Modeling




Most of the time  
there is a mix →  
Meet-in-the-Middle


# There are Many Languages/Notations



# Summary of Symbols Used in E-R Notation




entity set



relationship set



identifying  
relationship set  
for weak entity set



total participation  
of entity set in  
relationship

E
A1
A2
A2.1
A2.2
{A3}
A4()

attributes:  
simple (A1),  
composite (A2) and  
multivalued (A3)  
derived (A4)

E
<u>A1</u>


primary key

E
.....


discriminating  
attribute of  
weak entity set




# Symbols Used in E-R Notation (Cont.)




many-to-many  
relationship




one-to-one  
relationship




role  
indicator




many-to-one  
relationship




cardinality  
limits



ISA: generalization  
or specialization






total (disjoint)  
generalization



disjoint  
generalization

# Crow's Foot Notation

Many - to - One			
	Entity (with no attributes)		a one through many notation on one side of a relationship and a one and only one on the other
	Entity (with attributes field)		a zero through many notation on one side of a relationship and a one and only one on the other
	Entity (attributes field with columns)	 	a one through many notation on one side of a relationship and a zero or one notation on the other
Many-to-Many			
	Entity (attributes field with columns and variable number of rows)	 	a zero through many on both sides of a relationship
Relationships (Cardinality and Modality)			
	Zero or More	 	Many-to-Many
	One or More		a one through many on both sides of a relationship
	One and only One	 	a one and only one notation on one side of a relationship and a zero or one on the other
	Zero or One	 	a one and only one notation on both sides




# UML

- **UML**: Unified Modeling Language
- UML has many components to graphically model different aspects of an entire software system
- UML Class Diagrams correspond to E-R Diagram, but several differences.




# ER vs. UML Class Diagrams

## ER Diagram Notation




entity with attributes (simple, composite, multivalued, derived)




binary relationship



## Equivalent in UML




class with simple attributes and methods (attribute prefixes: + = public, - = private, # = protected)



relationship attributes



cardinality constraints




- \* Note reversal of position in cardinality constraint depiction




# ER vs. UML Class Diagrams


## ER Diagram Notation



n-ary relationships





overlapping  
generalization




disjoint  
generalization

## Equivalent in UML



overlapping



disjoint


- \* Generalization can use merged or separate arrows independent of disjoint/overlapping

# Modeling Summary

- If you are in a project/team meeting and draw a diagram on the whiteboard/chalk board, you are probably “visual modeling.”
- A visual notation is simply an agreement on precisely defining what the symbols, lines, colors, ... ... “mean.”
- Even within a well-defined notation, people are often lazy or class but use slightly different interpretations/notations.
- The two key concepts are:
  - Everyone in the process agrees on the meanings and adheres to them.
  - You do “just enough” modeling:
    - No modeling → You are randomly, frantically coding like a slam poet.
    - Too much modeling → You spend all of your time arguing and never deliver the project.
- We will take an approach to Crow’s Foot Notation in the class, but will also examine the modeling language that comes with the book. *This is a good fit with relational databases and relational model concepts.*

# Our First Example



# Modeling Columbia Courses – Conceptual



# Crow's Foot Notation

Relationships have two indicators. These are shown on both sides of the line.

- The first one (often called multiplicity) refers to the *maximum* number of times that an instance of one entity can be associated with instances in the related entity. It can be one or many.




- The second describes the *minimum* number of times one instance can be related to others. It can be zero or one, and accordingly describes the relationship as optional or mandatory.



<https://vertabelo.com/blog/crow-s-foot-notation/>

# Modeling Columbia Courses – Logical





You usually figure out if this is correct when you define in SQL and test use cases.

# Some Observations

- Foreign keys can be confusing and seem to sometimes be on the wrong end.
- Mandatory on both ends may require two new entities. The SQL INSERT statement operates on one table making INSERT complex. The data is temporarily inconsistent/invalid.
- Many-to-Many in SQL requires a 3<sup>rd</sup> entity called an Associative Entity.

For Example  
Mandatory on this end  
implies  
Not-NULL, FK on this end.



# Keys

# What are these Keys of which You Speak?



# Keys

- Let  $K \subseteq R$
- $K$  is a **superkey** of  $R$  if values for  $K$  are sufficient to identify a unique tuple of each possible relation  $r(R)$ 
  - Example:  $\{ID\}$  and  $\{ID, name\}$  are both superkeys of *instructor*.
- Superkey  $K$  is a **candidate key** if  $K$  is minimal  
Example:  $\{ID\}$  is a candidate key for *Instructor*
- One of the candidate keys is selected to be the **primary key**.
  - which one?
- **Foreign key** constraint: Value in one relation must appear in another
  - **Referencing** relation
  - **Referenced** relation
  - Example: *dept\_name* in *instructor* is a foreign key from *instructor* referencing *department*

# Our Model

- Student has three candidate keys:
  - UNI
  - CUID
  - email
- Faculty has two candidate keys:
  - UNI
  - email
- Section has two candidate keys:
  - callno
  - The combination of (dept\_code, faculty\_code, course\_no, section\_no, semester, year)
- Some types of key:
  - UNI, CUID, callno are *surrogate keys* as opposed to *natural keys*.
  - email and (dept\_code, faculty\_code, course\_no, section\_no, semester, year) are natural keys,

# Types of/Patterns for Keys

Key Type	Definition	Example	Primary Keys are Recommended to be immutable.
<b>Primary Key</b>	Uniquely identifies a record, no duplicates.	Employee ID, Student Roll Number	
<b>Super Key</b>	A combination of attributes that can uniquely identify records.	Employee ID + Name	
<b>Candidate Key</b>	A minimal super key that can act as a primary key.	Employee ID, Email Address	
<b>Alternate Key</b>	A candidate key is not chosen as the primary key.	Email Address (if Employee ID is primary)	
<b>Foreign Key</b>	Links tables by referencing a primary key from another table.	Department ID in Employee Table	
<b>Composite Key</b>	Multiple columns are used together to identify records uniquely.	Course ID + Student ID	
<b>Unique Key</b>	Ensures uniqueness in a column but allows one null value.	Email Address, Phone Number	

You may also hear the term *compound key* → Each of the attributes is a foreign key.

# Relational Model Relational Algebra



# Relational Algebra

- A procedural (**? Declarative**) language consisting of a set of operations that take one or two relations as input and produce a new relation as their result.

- Six basic operators

- select:  $\sigma$
- project:  $\Pi$
- union:  $\cup$
- set difference:  $-$
- Cartesian product:  $\times$
- JOIN ( $\bowtie$ )
- rename:  $\rho$

Covered last week.

- Let's do some practice with the dreaded RelaX calculator.
- Note that the RelaX calculator has an extended set of noncore operators.

# Let's Watch Professor Ferguson Do Examples

- Let's do some practice with the dreaded RelaX calculator.
- Pay attention to see how I use the “out of data” dataset associated with the recommended textbook.
- Note that the RelaX calculator has an extended set of noncore operators.
- Get the names of the *departments* with a budget  $\geq 100000$  or in the building ‘Painter’  
$$\pi \text{dept\_name} \\ (\sigma \text{budget} \geq 100000 \vee \text{building} = \text{'Painter'}) \text{ (department)}$$
Notes:
  - Select comparators from the menu, e.g.  $\geq$ .
  - Single quote matters.
  - Unlike SQL, strings are case sensitive.
  - Formally, relational algebra use subscript but this is hard to type.
- Rename attributes and operate on attributes in  $\pi$   
$$\pi \text{the\_course\_id} \leftarrow \text{course\_id}, \\ \text{sec\_id}, \text{year}, \text{something\_weird} \leftarrow \text{sec\_id} + \text{year} \\ \text{(section)}$$

# The Dreaded Relax Calculator

- Let's look at an online tool that you will use.
- RelaX (<https://dbis-uibk.github.io/relax/calc/local/uibk/local/0>)
- The calculator:
  - Has an older version of the data from the recommended textbook.  
(<https://dbis-uibk.github.io/relax/calc/gist/4f7866c17624ca9dfa85ed2482078be8/relax-silberschatz-english.txt/0>)
  - You can also upload new data.
- Some queries:
  - $\sigma \text{dept\_name}='Comp. Sci.' \vee \text{dept\_name}='History'$  (department)
  - $\pi \text{name}, \text{dept\_name}$  (instructor)
  - $\pi \text{ID}, \text{name}$   
 $\quad \sigma \text{dept\_name}='Comp. Sci.'$  (instructor)  
)



# Cartesian-Product Operation

- The Cartesian-product operation (denoted by  $\times$ ) allows us to combine information from any two relations.
- Example: the Cartesian product of the relations *instructor* and *teaches* is written as:  
$$\textit{instructor} \times \textit{teaches}$$
- We construct a tuple of the result out of each possible pair of tuples: one from the *instructor* relation and one from the *teaches* relation (see next slide)
- Since the instructor *ID* appears in both relations we distinguish between these attribute by attaching to the attribute the name of the relation from which the attribute originally came.
  - $\textit{instructor.ID}$
  - $\textit{teaches.ID}$

$$R \times R \rightarrow R \quad f(x,y) \rightarrow z$$



# Composition of Relational Operations

- The result of a relational-algebra operation is relation and therefore of relational-algebra operations can be composed together into a **relational-algebra expression**.
- Consider the query -- Find the names of all instructors in the Physics department.

$$\Pi_{name}(\sigma_{dept\_name = "Physics"}(instructor))$$

- Instead of giving the name of a relation as the argument of the projection operation, we give an expression that evaluates to a relation.




# The *instructor* x *teaches* table

- This only sort of makes sense. The result is:
  - Every possible combination of the form (instructor, teaches)
  - Even if the instructor is NOT the instructor in the teaches row.
- Examining in SQL makes a little clearer.
  - We will see in later in this lecture.
  - Confusingly, in SQL cartesian Product is **JOIN** but there is also a relational **JOIN** operator.

Instructor.ID	name	dept_name	salary	teaches.ID	course_id	sec_id	semester	year
10101	Srinivasan	Comp. Sci.	65000	10101	CS-101	1	Fall	2017
10101	Srinivasan	Comp. Sci.	65000	10101	CS-315	1	Spring	2018
10101	Srinivasan	Comp. Sci.	65000	10101	CS-347	1	Fall	2017
10101	Srinivasan	Comp. Sci.	65000	12121	FIN-201	1	Spring	2018
10101	Srinivasan	Comp. Sci.	65000	15151	MU-199	1	Spring	2018
10101	Srinivasan	Comp. Sci.	65000	22222	PHY-101	1	Fall	2017
...	...	...	...	...	...	...	...	...
...	...	...	...	...	...	...	...	...
12121	Wu	Finance	90000	10101	CS-101	1	Fall	2017
12121	Wu	Finance	90000	10101	CS-315	1	Spring	2018
12121	Wu	Finance	90000	10101	CS-347	1	Fall	2017
12121	Wu	Finance	90000	12121	FIN-201	1	Spring	2018
12121	Wu	Finance	90000	15151	MU-199	1	Spring	2018
12121	Wu	Finance	90000	22222	PHY-101	1	Fall	2017
...	...	...	...	...	...	...	...	...
...	...	...	...	...	...	...	...	...
15151	Mozart	Music	40000	10101	CS-101	1	Fall	2017
15151	Mozart	Music	40000	10101	CS-315	1	Spring	2018
15151	Mozart	Music	40000	10101	CS-347	1	Fall	2017
15151	Mozart	Music	40000	12121	FIN-201	1	Spring	2018
15151	Mozart	Music	40000	15151	MU-199	1	Spring	2018
15151	Mozart	Music	40000	22222	PHY-101	1	Fall	2017
...	...	...	...	...	...	...	...	...
...	...	...	...	...	...	...	...	...
22222	Einstein	Physics	95000	10101	CS-101	1	Fall	2017
22222	Einstein	Physics	95000	10101	CS-315	1	Spring	2018
22222	Einstein	Physics	95000	10101	CS-347	1	Fall	2017
22222	Einstein	Physics	95000	12121	FIN-201	1	Spring	2018
22222	Einstein	Physics	95000	15151	MU-199	1	Spring	2018
22222	Einstein	Physics	95000	22222	PHY-101	1	Fall	2017
...	...	...	...	...	...	...	...	...
...	...	...	...	...	...	...	...	...

# Simpler Example

- Assume we have two tables
  - S has two columns, 5 rows.
  - T has two columns, 4 rows.
- $S \times T$  has
  - 4 columns.
  - 20 rows.
- Cartesian product does not come up a lot in applications.
- There are cases in optimization and other scenarios in which:
  - You want to generate all possible combinations.
  - Score, rate, rank etc. to determine best choices.





# Join Operation

- The Cartesian-Product

*instructor X teaches*

associates every tuple of instructor with every tuple of teaches.

- Most of the resulting rows have information about instructors who did NOT teach a particular course.

- To get only those tuples of “*instructor X teaches*” that pertain to instructors and the courses that they taught, we write:

$\sigma_{instructor.id = teaches.id} (instructor \times teaches)$ )

- We get only those tuples of “*instructor X teaches*” that pertain to instructors and the courses that they taught.
- The result of this expression, shown in the next slide

A fundamental definition:

- $\sigma_{instructor.ID=teaches.ID} (instructor \times teaches) = instructor \bowtie teaches$
- $\bowtie$  is the JOIN operations.



## Join Operation (Cont.)

- The **join** operation allows us to combine a select operation and a Cartesian-Product operation into a single operation.
- Consider relations  $r (R)$  and  $s (S)$
- Let “theta” be a predicate on attributes in the schema R “union” S. The join operation  $r \bowtie_{\theta} s$  is defined as follows:

$$r \bowtie_{\theta} s = \sigma_{\theta} (r \times s)$$

- Thus

$$\sigma_{instructor.id = teaches.id} (instructor \times teaches))$$

- Can equivalently be written as

*instructor*  $\bowtie$  *Instructor.id = teaches.id teaches.*

# Let's Watch Professor Ferguson Do Examples

- More fun with RelaX
- Validate the definition with an example
  - $\sigma \text{instructor}.\text{dept\_name}=\text{department}.\text{dept\_name}$  ( $\text{instructor} \times \text{department}$ )
  - $\text{instructor} \bowtie \text{department}$
  - $\text{instructor} \bowtie \sigma \text{instructor}.\text{dept\_name}=\text{department}.\text{dept\_name} \text{ department}$
  - Note: You sometimes have to fully qualify attribute names.

- Let's pull it all together ... ...

```
 $\sigma \text{instructor\_department}='Comp. Sci.'$ 
```

```
(
```

```
 $\pi$ 
```

```
     $\text{instructor\_id} \leftarrow \text{instructor.ID}$ ,  $\text{instructor\_name} \leftarrow \text{instructor.name}$ ,  
     $\text{instructor\_department} \leftarrow \text{instructor.dept\_name}$ ,  $s\_id$   
 $(\text{instructor} \bowtie \text{instructor.ID}=i\_id \text{ advisor})$ 
```

```
 $\bowtie \text{student\_id}=s\_id$ 
```

```
( $\pi$ 
```

```
     $\text{student\_id} \leftarrow \text{student.ID}$ ,  $\text{student\_name} \leftarrow \text{student.name}$ ,  
 $\text{student\_department} \leftarrow \text{student.dept\_name}$  ( $\text{student}$ )
```

```
)
```

## Notes:

- Show them how to break this down Barney Style.
- Show them approaches to testing and verifying.
- Being able to verify is CRITICAL if you use generative AI.

# SQL



# SQL Parts

- DML -- provides the ability to query information from the database and to insert tuples into, delete tuples from, and modify tuples in the database.
- integrity – the DDL includes commands for specifying integrity constraints.
- View definition -- The DDL includes commands for defining views.
- Transaction control –includes commands for specifying the beginning and ending of transactions.
- Embedded SQL and dynamic SQL -- define how SQL statements can be embedded within general-purpose programming languages.
- Authorization – includes commands for specifying access rights to relations and views.

# SQL Language Statements

The core SQL language statements are:

- SELECT: Implements both  $\sigma$ ,  $\pi$
- INSERT
- UPDATE
- DELETE
- CREATE TABLE
- ALTER TABLE
- JOIN, which is an operator within SELECT.
- Many, if not most, SQL statements:
  - Implement multiple relational algebra expressions.
  - Cannot easily (or at all) be represented in relational algebra.

$$\begin{aligned} \pi \text{ ID, name} ( \\ & \quad \sigma \text{ dept\_name='Comp. Sci.'} (\text{instructor}) \\ ) \\ = \\ \text{SELECT ID, name FROM instructor} \\ \text{WHERE} \\ \text{dept\_name='Comp. Sci.'} \end{aligned}$$



# Basic Query Structure

- A typical SQL query has the form:

```
select A1, A2, ..., An
from r1, r2, ..., rm
where P
```

- $A_i$  represents an attribute
- $R_i$  represents a relation
- $P$  is a predicate.
- The result of an SQL query is a relation.

Multiple tables and cause  
confusion and errors because  
you are doing a cross=product.  
Be careful

## Note:

- The SELECT ... FROM ... WHERE ... Combines two relational operators,  $\sigma$  and  $\Pi$ .
- Actually, it also combines other operators, e.g.  $\times$



# The select Clause

- The **select** clause lists the attributes desired in the result of a query
  - corresponds to the projection operation of the relational algebra
- Example: find the names of all instructors:

```
select name  
      from instructor
```
- NOTE: SQL names are case insensitive (i.e., you may use upper- or lower-case letters.)
  - E.g., *Name*  $\equiv$  *NAME*  $\equiv$  *name*
  - Some people use upper case wherever we use bold font.



# The select Clause (Cont.)

- SQL allows duplicates in relations as well as in query results.
- To force the elimination of duplicates, insert the keyword **distinct** after `select`.
- Find the department names of all instructors, and remove duplicates
  - ```
select distinct dept_name  
      from instructor
```
- The keyword **all** specifies that duplicates should not be removed.

```
select all dept_name  
      from instructor
```



# The select Clause (Cont.)

- An asterisk in the select clause denotes “all attributes”

```
select *
from instructor
```

- An attribute can be a literal with no **from** clause

```
select '437'
```

- Results is a table with one column and a single row with value “437”
- Can give the column a name using:

```
select '437' as FOO
```

- An attribute can be a literal with **from** clause

```
select 'A'
from instructor
```

- Result is a table with one column and  $N$  rows (number of tuples in the *instructors* table), each row with value “A”



## The select Clause (Cont.)

- The **select** clause can contain arithmetic expressions involving the operation, +, -, \*, and /, and operating on constants or attributes of tuples.
  - The query:

```
select ID, name, salary/12  
from instructor
```

would return a relation that is the same as the *instructor* relation, except that the value of the attribute *salary* is divided by 12.

- Can rename “salary/12” using the **as** clause:

```
select ID, name, salary/12 as monthly_salary
```



# The where Clause

- The **where** clause specifies conditions that the result must satisfy
  - Corresponds to the selection predicate of the relational algebra.
- To find all instructors in Comp. Sci. dept

```
select name  
from instructor  
where dept_name = 'Comp. Sci.'
```

- SQL allows the use of the logical connectives **and**, **or**, and **not**
- The operands of the logical connectives can be expressions involving the comparison operators <, <=, >, >=, =, and <>.
- Comparisons can be applied to results of arithmetic expressions
- To find all instructors in Comp. Sci. dept with salary > 80000

```
select name  
from instructor  
where dept_name = 'Comp. Sci.' and salary > 80000
```



# The from Clause

- The **from** clause lists the relations involved in the query
  - Corresponds to the Cartesian product operation of the relational algebra.
- Find the Cartesian product *instructor X teaches*

```
select *
  from instructor, teaches
```

  - generates every possible instructor – teaches pair, with all attributes from both relations.
  - For common attributes (e.g., *ID*), the attributes in the resulting table are renamed using the relation name (e.g., *instructor.ID*)
- Cartesian product not very useful directly, but useful combined with where-clause condition (selection operation in relational algebra).



# Examples

- Find the names of all instructors who have taught some course and the course\_id
  - **select name, course\_id  
from instructor , teaches  
where instructor.ID = teaches.ID**
- Find the names of all instructors in the Art department who have taught some course and the course\_id
  - **select name, course\_id  
from instructor , teaches  
where instructor.ID = teaches.ID **and** instructor.dept\_name = 'Art'**



# Data Definition Language

The SQL data-definition language (DDL) allows the specification of information about relations, including:

- The schema for each relation.
- The type of values associated with each attribute.
- The Integrity constraints
- The set of indices to be maintained for each relation.
- Security and authorization information for each relation.
- The physical storage structure of each relation on disk.



# Create Table Construct

- An SQL relation is defined using the **create table** command:

**create table** *r*

( $A_1 D_1, A_2 D_2, \dots, A_n D_n,$   
(integrity-constraint<sub>1</sub>),  
...,  
(integrity-constraint<sub>k</sub>))

- r* is the name of the relation
  - each  $A_i$  is an attribute name in the schema of relation *r*
  - $D_i$  is the data type of values in the domain of attribute  $A_i$
- Example:

```
create table instructor (  
    ID          char(5),  
    name        varchar(20),  
    dept_name   varchar(20),  
    salary      numeric(8,2))
```



# The Rename Operation

- The SQL allows renaming relations and attributes using the **as** clause:

*old-name as new-name*

- Find the names of all instructors who have a higher salary than some instructor in 'Comp. Sci'.

- **select distinct** *T.name*  
**from** *instructor as T, instructor as S*  
**where** *T.salary > S.salary and S.dept\_name = 'Comp. Sci.'*

This is an example of using  
Cartesian Product.

- Keyword **as** is optional and may be omitted

*instructor as T*  $\equiv$  *instructor T*



# Null Values

- It is possible for tuples to have a null value, denoted by **null**, for some of their attributes
- **null** signifies an **unknown value** or that a **value does not exist**.
- The result of any arithmetic expression involving **null** is **null**
  - Example:  $5 + \text{null}$  returns **null**
- The predicate **is null** can be used to check for null values.
  - Example: Find all instructors whose salary is null.

```
select name  
from instructor  
where salary is null
```

- The predicate **is not null** succeeds if the value on which it is applied is not null.

## Note:

- **NULL is an extremely important concept.**
- **You will find it hard to understand for a while.**



## Null Values (Cont.)

- SQL treats as **unknown** the result of any comparison involving a null value (other than predicates **is null** and **is not null**).
  - Example:  $5 < \text{null}$  or  $\text{null} < > \text{null}$  or  $\text{null} = \text{null}$
- The predicate in a **where** clause can involve Boolean operations (**and**, **or**, **not**); thus the definitions of the Boolean operations need to be extended to deal with the value **unknown**.
  - **and** :  $(\text{true and unknown}) = \text{unknown}$ ,  
 $(\text{false and unknown}) = \text{false}$ ,  
 $(\text{unknown and unknown}) = \text{unknown}$
  - **or**:  $(\text{unknown or true}) = \text{true}$ ,  
 $(\text{unknown or false}) = \text{unknown}$   
 $(\text{unknown or unknown}) = \text{unknown}$
- Result of **where** clause predicate is treated as *false* if it evaluates to *unknown*



# Modification of the Database

- Deletion of tuples from a given relation.
- Insertion of new tuples into a given relation
- Updating of values in some tuples in a given relation



# Deletion

- Delete all instructors

**delete from** *instructor*

- Delete all instructors from the Finance department

**delete from** *instructor*  
**where** *dept\_name*= 'Finance';

- *Delete all tuples in the instructor relation for those instructors associated with a department located in the Watson building.*

**delete from** *instructor*  
**where** *dept\_name* **in** (**select** *dept\_name*  
         **from** *department*  
         **where** *building* = 'Watson');



## Deletion (Cont.)

- Delete all instructors whose salary is less than the average salary of instructors

```
delete from instructor  
where salary < (select avg (salary)  
          from instructor);
```

- Problem: as we delete tuples from *instructor*, the average salary changes
- Solution used in SQL:
  1. First, compute **avg** (*salary*) and find all tuples to delete
  2. Next, delete all tuples found above (without recomputing **avg** or retesting the tuples)



# Insertion

- Add a new tuple to *course*

```
insert into course
values ('CS-437', 'Database Systems', 'Comp. Sci.', 4);
```

- or equivalently

```
insert into course (course_id, title, dept_name, credits)
values ('CS-437', 'Database Systems', 'Comp. Sci.', 4);
```

- Add a new tuple to *student* with *tot\_creds* set to null

```
insert into student
values ('3003', 'Green', 'Finance', null);
```



## Insertion (Cont.)

- Make each student in the Music department who has earned more than 144 credit hours an instructor in the Music department with a salary of \$18,000.

```
insert into instructor
    select ID, name, dept_name, 18000
        from student
    where dept_name = 'Music' and total_cred > 144;
```

- The **select from where** statement is evaluated fully before any of its results are inserted into the relation.

Otherwise queries like

```
insert into table1 select * from table1
```

would cause problem



# Updates

- Give a 5% salary raise to all instructors

```
update instructor  
    set salary = salary * 1.05
```

- Give a 5% salary raise to those instructors who earn less than 70000

```
update instructor  
    set salary = salary * 1.05  
    where salary < 70000;
```

- Give a 5% salary raise to instructors whose salary is less than average

```
update instructor  
    set salary = salary * 1.05  
    where salary < (select avg (salary)  
                    from instructor);
```



# Updates (Cont.)

- Increase salaries of instructors whose salary is over \$100,000 by 3%, and all others by a 5%
  - Write two **update** statements:

```
update instructor
    set salary = salary * 1.03
    where salary > 100000;
update instructor
    set salary = salary * 1.05
    where salary <= 100000;
```

- The order is important
- Can be done better using the **case** statement (next slide)



# Case Statement for Conditional Updates

- Same query as before but with case statement

```
update instructor  
set salary = case  
    when salary <= 100000 then salary * 1.05  
    else salary * 1.03  
end
```



# Joined Relations

- **Join operations** take two relations and return as a result another relation.
- A join operation is a Cartesian product which requires that tuples in the two relations match (under some condition). It also specifies the attributes that are present in the result of the join
- The join operations are typically used as subquery expressions in the **from** clause
- Three types of joins:
  - Natural join
  - Inner join
  - Outer join

## Notes:

- You will also hear terms like equi-join, non-equi-join, theta join, semi-join, ... ....
- I ask for definitions on exams, but you can just look them up.



# Natural Join in SQL

- Natural join matches tuples with the same values for all common attributes, and retains only one copy of each common column.
- List the names of instructors along with the course ID of the courses that they taught
  - **`select name, course_id  
from students, takes,  
where student.ID = takes.ID;`**
- Same query in SQL with “natural join” construct
  - **`select name, course_id  
from student natural join takes;`**



## Natural Join in SQL (Cont.)

- The **from** clause can have multiple relations combined using natural join:

```
select A1, A2, ... An  
from r1 natural join r2 natural join .. natural join rn  
where P;
```



# Student Relation

| <i>ID</i> | <i>name</i> | <i>dept_name</i> | <i>tot_cred</i> |
|-----------|-------------|------------------|-----------------|
| 00128     | Zhang       | Comp. Sci.       | 102             |
| 12345     | Shankar     | Comp. Sci.       | 32              |
| 19991     | Brandt      | History          | 80              |
| 23121     | Chavez      | Finance          | 110             |
| 44553     | Peltier     | Physics          | 56              |
| 45678     | Levy        | Physics          | 46              |
| 54321     | Williams    | Comp. Sci.       | 54              |
| 55739     | Sanchez     | Music            | 38              |
| 70557     | Snow        | Physics          | 0               |
| 76543     | Brown       | Comp. Sci.       | 58              |
| 76653     | Aoi         | Elec. Eng.       | 60              |
| 98765     | Bourikas    | Elec. Eng.       | 98              |
| 98988     | Tanaka      | Biology          | 120             |



# Takes Relation

| <i>ID</i> | <i>course_id</i> | <i>sec_id</i> | <i>semester</i> | <i>year</i> | <i>grade</i> |
|-----------|------------------|---------------|-----------------|-------------|--------------|
| 00128     | CS-101           | 1             | Fall            | 2017        | A            |
| 00128     | CS-347           | 1             | Fall            | 2017        | A-           |
| 12345     | CS-101           | 1             | Fall            | 2017        | C            |
| 12345     | CS-190           | 2             | Spring          | 2017        | A            |
| 12345     | CS-315           | 1             | Spring          | 2018        | A            |
| 12345     | CS-347           | 1             | Fall            | 2017        | A            |
| 19991     | HIS-351          | 1             | Spring          | 2018        | B            |
| 23121     | FIN-201          | 1             | Spring          | 2018        | C+           |
| 44553     | PHY-101          | 1             | Fall            | 2017        | B-           |
| 45678     | CS-101           | 1             | Fall            | 2017        | F            |
| 45678     | CS-101           | 1             | Spring          | 2018        | B+           |
| 45678     | CS-319           | 1             | Spring          | 2018        | B            |
| 54321     | CS-101           | 1             | Fall            | 2017        | A-           |
| 54321     | CS-190           | 2             | Spring          | 2017        | B+           |
| 55739     | MU-199           | 1             | Spring          | 2018        | A-           |
| 76543     | CS-101           | 1             | Fall            | 2017        | A            |
| 76543     | CS-319           | 2             | Spring          | 2018        | A            |
| 76653     | EE-181           | 1             | Spring          | 2017        | C            |
| 98765     | CS-101           | 1             | Fall            | 2017        | C-           |
| 98765     | CS-315           | 1             | Spring          | 2018        | B            |
| 98988     | BIO-101          | 1             | Summer          | 2017        | A            |
| 98988     | BIO-301          | 1             | Summer          | 2018        | <i>null</i>  |



## student natural join takes

| <i>ID</i> | <i>name</i> | <i>dept_name</i> | <i>tot_cred</i> | <i>course_id</i> | <i>sec_id</i> | <i>semester</i> | <i>year</i> | <i>grade</i> |
|-----------|-------------|------------------|-----------------|------------------|---------------|-----------------|-------------|--------------|
| 00128     | Zhang       | Comp. Sci.       | 102             | CS-101           | 1             | Fall            | 2017        | A            |
| 00128     | Zhang       | Comp. Sci.       | 102             | CS-347           | 1             | Fall            | 2017        | A-           |
| 12345     | Shankar     | Comp. Sci.       | 32              | CS-101           | 1             | Fall            | 2017        | C            |
| 12345     | Shankar     | Comp. Sci.       | 32              | CS-190           | 2             | Spring          | 2017        | A            |
| 12345     | Shankar     | Comp. Sci.       | 32              | CS-315           | 1             | Spring          | 2018        | A            |
| 12345     | Shankar     | Comp. Sci.       | 32              | CS-347           | 1             | Fall            | 2017        | A            |
| 19991     | Brandt      | History          | 80              | HIS-351          | 1             | Spring          | 2018        | B            |
| 23121     | Chavez      | Finance          | 110             | FIN-201          | 1             | Spring          | 2018        | C+           |
| 44553     | Peltier     | Physics          | 56              | PHY-101          | 1             | Fall            | 2017        | B-           |
| 45678     | Levy        | Physics          | 46              | CS-101           | 1             | Fall            | 2017        | F            |
| 45678     | Levy        | Physics          | 46              | CS-101           | 1             | Spring          | 2018        | B+           |
| 45678     | Levy        | Physics          | 46              | CS-319           | 1             | Spring          | 2018        | B            |
| 54321     | Williams    | Comp. Sci.       | 54              | CS-101           | 1             | Fall            | 2017        | A-           |
| 54321     | Williams    | Comp. Sci.       | 54              | CS-190           | 2             | Spring          | 2017        | B+           |
| 55739     | Sanchez     | Music            | 38              | MU-199           | 1             | Spring          | 2018        | A-           |
| 76543     | Brown       | Comp. Sci.       | 58              | CS-101           | 1             | Fall            | 2017        | A            |
| 76543     | Brown       | Comp. Sci.       | 58              | CS-319           | 2             | Spring          | 2018        | A            |
| 76653     | Aoi         | Elec. Eng.       | 60              | EE-181           | 1             | Spring          | 2017        | C            |
| 98765     | Bourikas    | Elec. Eng.       | 98              | CS-101           | 1             | Fall            | 2017        | C-           |
| 98765     | Bourikas    | Elec. Eng.       | 98              | CS-315           | 1             | Spring          | 2018        | B            |
| 98988     | Tanaka      | Biology          | 120             | BIO-101          | 1             | Summer          | 2017        | A            |
| 98988     | Tanaka      | Biology          | 120             | BIO-301          | 1             | Summer          | 2018        | <i>null</i>  |



# Dangerous in Natural Join

- Beware of unrelated attributes with same name which get equated incorrectly
- Example -- List the names of students instructors along with the titles of courses that they have taken
  - Correct version

```
select name, title  
from student natural join takes, course  
where takes.course_id = course.course_id;
```

- Incorrect version
- ```
select name, title  
from student natural join takes natural join course;
```
- This query omits all (student name, course title) pairs where the student takes a course in a department other than the student's own department.
  - The correct version (above), correctly outputs such pairs.



# Natural Join with Using Clause

- To avoid the danger of equating attributes erroneously, we can use the “**using**” construct that allows us to specify exactly which columns should be equated.
- Query example

```
select name, title  
from (student natural join takes) join course using (course_id)
```



# Join Condition

- The **on** condition allows a general predicate over the relations being joined
- This predicate is written like a **where** clause predicate except for the use of the keyword **on**
- Query example

```
select *  
from student join takes on student_ID = takes_ID
```

- The **on** condition above specifies that a tuple from *student* matches a tuple from *takes* if their *ID* values are equal.
- Equivalent to:

```
select *  
from student , takes  
where student_ID = takes_ID
```



## Join Condition (Cont.)

- The **on** condition allows a general predicate over the relations being joined.
- This predicate is written like a **where** clause predicate except for the use of the keyword **on**.
- Query example

```
select *
from student join takes on student_ID = takes_ID
```

- The **on** condition above specifies that a tuple from *student* matches a tuple from *takes* if their *ID* values are equal.
- Equivalent to:

```
select *
from student, takes
where student_ID = takes_ID
```